


Feb 24th – 26th, 2026




PPHPS – GPU Programming with OpenMP Target Offloading

Sebastian Kuckuk, *Erlangen National High Performance Computing Center (NHR@FAU)*

GPUs in HPC – Motivation

- Promises
 - Massive parallelism and performance
 - Good performance in relation to energy (FLOPs per Watt)
- Already widespread in the HPC landscape
 - c.f. Top500 list (<https://www.top500.org/lists/top500/2025/11/>)
 - 9 out of the top 10 supercomputers are equipped with GPUs
 - 4 NVIDIA (2 x GH200, H100, A100)
 - 4 AMD (MI300A, 3 x MI250X)
 - 1 Intel (GPU Max Series)
- Where does the performance come from?




Helma is #58
SuperMUC-NG is #91
SuperMUC-NG Ph. 2 is #108
Alex is #322
Fritz is #350

GPU Performance

- If GPUs are mainly about performance – what is performance?
- Two primary factors
 - How fast can the meaningful computation be done?
Usually given as computational throughput, e.g. FLOP/s
 - How fast can the data be transferred to where the computation is happening (and back)?
Usually given as sustained bandwidth, e.g. GB/s
- What are GPUs doing differently?

GPU vs CPU


- CPUs have few (~100) powerful cores
- GPUs have many simplistic 'cores' (10 000s)



- Organized in sockets




- Organized in 100s of streaming multiprocessors (SMs)




GPU vs CPU

- Both CPU and GPU have a distinct main memory (for most architectures)





DRAM (100s GB)



DRAM (1000s GB)


GPU vs CPU

- The memory of GPUs
 - Is optimized for bandwidth
 - Has a high latency
- The memory of CPUs
 - Is optimized for latency
 - Has a high capacity



GPU vs CPU

- Both CPU and GPU have a distinct main memory (for most architectures)
- They communicate via slow interconnects (for most architectures)



Introduction – CPU vs GPU

- Why not use GPUs exclusively?
 - Let's consider a simple vector copy benchmark to assess sustained bandwidths
- One A100 40GB GPU (Alex) vs one Sapphire Rapids node (Fritz)
 - ~ 1.3 TB/s vs ~ 260 GB/s => 5x faster
- But: One SM of one A100: ~ 90 GB/s => 3x slower
- Serial execution
 - ~ 0.3 GB/s vs ~ 20 GB/s => 67x slower

Introduction – CPU vs GPU

- Why not use GPUs exclusively?
- GPUs deal best with
 - Massive parallelism
 - at least 10 000s of threads to saturate computation*
 - at least 100 000s of threads to saturate memory*
 - Structured computations
 - Ideally each thread does the same operation but on different parts of a structured data set
- CPUs deal much better with unstructured, low-parallelism computations

GPU Programming Approaches

Options are plentiful

1. Avoid GPU programming all together
 - GPU-accelerated libraries
2. Let the compiler do the job
 - Modern C++
 - Pragma-based approaches
(OpenMP, OpenACC)
3. Get your hands dirty and do the technical work
 - CUDA, HIP, oneAPI, ...
4. Do all the work – but now with added performance portability
 - Software layers (Kokkos, ...)

GPU Programming Approaches

Options are plentiful

1. Avoid GPU programming all together
 - GPU-accelerated libraries
2. Let the compiler do the job
 - Modern C++
 - **Pragma-based approaches**
(OpenMP, OpenACC)
3. Get your hands dirty and do the technical work
 - CUDA, HIP, oneAPI, ...
4. Do all the work – but now with added performance portability
 - Software layers (Kokkos, ...)